

Reticular Observer Architectures for Governable AI-Assisted Work

A Reticular Framework for Epistemic Abstraction, Computability, and Proof-Aware Systems

Gianluca Conte

Independent Researcher

contegianluca@hotmail.com

May 2026

Abstract

This paper breathes with two lungs. The first is constructive: it argues that semantics is computable in an observer-relative, horizon-bounded sense, and that this makes possible a different way of building knowledge systems — mining a scientific domain with AI into a typed, computable reticulum, and answering by navigating that reticulum rather than regenerating answers. The destination is a governed oracle: a system that returns persistent, end-to-end traceable answers, with abstention and contradiction as first-class outcomes. The construction is not merely proposed; the author applied the framework to itself, using a pipeline of specialised AI agents to build the reference architecture that witnesses it — with an explicit caveat about what that does and does not prove. The second lung is disciplinary: the same construction makes the characteristic failure of AI, the undeclared approximation or hallucination, visible and governable, by treating every promotion of a pattern into an object as a typed, auditable reification carrying epistemic debt that blocks downstream authoritative use until discharged.

Reticular Local Abstraction (RLA) supplies the grammar of levels, transmissions and horizons; Compact Reticular Computability (CRC) the computability discipline under a declared horizon; epistemic convolution (ECNN) the operational mechanism; reticular observer architectures (ROA) the governance compression in which reifications, debts, proof states and terminal states become first-class. A worked case — a fully computable reticulum for a generalist bryophyte, assembled from the published literature — demonstrates the central method, science as node. The framework reuses established results — Rice’s theorem, provenance models, selective prediction, information-flow discipline — openly; its contribution is their composition into a single governable observer. A reference architecture (A-OSP) is an implementation witness gated by a narrow success proxy, not empirical validation. A result on the propagation of undecidability does real work, defeating the intuition of compositional decidability, and the paper states its own falsification conditions. The paper does not claim a completed mathematical theory, empirical validation, production readiness, or a theory of artificial consciousness.

Keywords: computable semantics, science as node, epistemic twin, AI knowledge mining, hallucination as declared debt, governed oracle, controlled reification, epistemic debt propagation, mandatory abstention, blocking debt, self-application, selective prediction, provenance, proof-aware systems.

PART I — THE CONSTRUCTION: WHAT CAN BE BUILT

1. The Problem, and a Different Shape of Answer

The bottleneck in AI-assisted work is no longer fluent generation; it is reconstruction, provenance, proof and governance. A prompt-first workflow generates a convincing report, analysis or compliance artefact, but afterwards an organisation often cannot say which evidence supported each claim, what was inferred rather than supported, what was missing, or whether the result was approved, merely reviewed, or only drafted. The dominant pattern prompt → answer hides the observer behind the answer.

The observer-compiler picture is different. Many of the most useful outputs of AI work are not final answers but intermediate, typed artefacts — definitions, taxonomies, blueprints, issue chains, graph fragments, challenge suites, review checklists — which a human curator corrects, compresses, validates and reuses as input for later steps. Cumulative intelligence is not located inside a single model call; it is distributed across calls, human decisions, files, versions, traces and reviews.

A prompt, on this reading, is the design of a natural-language operator over a field, not a request for a sentence. This paper proposes a shape that makes that structure explicit: bounded domain material → explicit observer structure → typed epistemic artefacts → validation state → proof / witness / review / governance. A reticular observer architecture is a finite, horizon-relative structure specifying what can be represented, how distinctions are transmitted or collapsed, which artefacts can be produced, and under which validation and governance conditions they may be used.

Framework



Instantiation — supplier payment control



Controlled reification and the resulting epistemic debt are an explicit, blocking transition — not an implicit side effect.

Figure 1. The reticular observer pipeline (top) and an instantiation on a compliance task (bottom). Controlled reification and the resulting epistemic debt are an explicit, blocking transition.

A single example runs through the paper, drawn from compliance review because that domain stresses evidence, its absence, accountability and human review. An organisation reviews whether supplier payments are adequately controlled; an analyst conducts interviews, and the transcripts

are the bounded domain material. The horizon admits questions about which controls exist, where evidence is missing and what follow-up is required, and excludes a final legal opinion on liability. The levels are raw answers, atomic facts, links between them, and a compact synthesis; the interviews yield, among others, the atom “a CFO-approval threshold exists” and the evidence gap “no formal exception-handling procedure was found.” The paper returns to this example at each layer — collapse, abstention, reification — so that the abstract machinery always has a concrete referent.

2. Computable Semantics: the Organizing Thesis

The framework has one spine. The thesis is that semantics is computable in an observer-relative, horizon-bounded sense: the observation and stabilization of meaning — selecting a domain, encoding it, transmitting distinctions across levels, collapsing detail, inducing labels, promoting patterns into manipulable objects, deciding terminal states — is an effective, bounded, horizon-relative process that emits typed epistemic artefacts. An answer is one artefact; unknown, contradiction, out-of-horizon and review-required are equally first-class, computable terminal states. A responsible abstention is a successful computation; a fluent but unsupported answer is a defective one.

What the thesis does and does not claim

Claimed: the disciplined construction and governance of an observer over a semantic field is an effective, horizon-relative process; semantic and artefactual fields become objects over which procedures operate. Not claimed: that meaning reduces to truth, that any judgement is correct because it was computed, or that the process is unbounded. It is always relative to a declared horizon H and an operational envelope.

3. Science as Node: Building the Horizon from Knowledge Itself

The central method treats a scientific domain as a node: not a backdrop, but the material from which the observer’s horizon and internal epistemology are constructed. One mines the published science of a domain — its levels, variables, equations, parameter ranges, feedbacks — with AI assistance, and harmonizes them into a compact reticulum that is internally autonomous (its explanatory inventory is inside the lattice), epistemically closed (it answers or responsibly fails within its horizon), and fully computable. The product is not the domain; it is an epistemic twin of it: a computable structure whose synthetic consequences can be compared against what the science asserts. The honest description of what AI contributes here is original epistemic artefact synthesis, not truth extraction: a model helps assemble candidate structures — taxonomies, reticula, hypotheses, gap analyses — whose originality lies in arrangement, correlation and operationalisation, never in a claim to have produced new truth.

This inverts the usual relation between a model and knowledge. The single source of truth is not a model’s weights, nor a chat transcript, but the governed scientific reticulum itself — science, mined by a human through AI, accumulated as a typed and traceable structure rather than re-generated at every call. The observer’s horizon H is read off the science, not imposed a priori; knowledge is composed into a navigable lattice. The mining field, in practice, is the whole working substrate — repository, source of truth, issues, blueprints, code, logs, graph — and the filters are prompts, agents, checklists and matrix constraints (see RLA grammar, §7).

4. The Bryophyte: a Worked Epistemic Twin

The method is demonstrated, not merely asserted. A fully computable RLA topology for a generalist bryophyte is assembled from the plant-physiology and ecology literature: twenty abstraction levels from genomic and epigenetic regulation, through hormone and reactive-oxygen dynamics, tissue mechanics and reproduction, up to microenvironment and macro-climate forcing; forty-two parameters with units, literature ranges and operational reconduction rules; families of update equations, threshold actions and rare-event injectors; collapse and emergence metrics; climate scenarios and an adaptive epigenetic-buffering mechanism. Every explanatory and predictive element is internal to the lattice, and the dynamics are Turing-computable. The bryophyte is treated as an epistemic model, not an ultimate ontology of life: its levels and equations encode disciplined, revisable scientific knowledge.

Epistemic indistinguishability — a falsifiable prediction


The framework makes a testable claim: synthetic trajectories generated by the compact reticulum should be epistemically indistinguishable, under blind expert inspection, from empirical observations under comparable conditions. Concretely: generate simulated trajectories of key observables; collect real trajectories from controlled microcosms; mix, anonymise, and ask a domain expert to classify each as real or synthetic. If discrimination is no better than a chosen threshold, the topology is empirically adequate relative to that observable set — a Turing test for computed versus measured science. This is a design goal, not yet a validated result; stating it as falsifiable is the point. (Full topology and protocol: Annex B.)

5. Mining, Hallucination, and the Undeclared Approximation

Mining science with AI exposes a precise and important failure mode, and the framework's treatment of it is one of its strongest practical moves. Suppose the miner needs a specific equation of the bryophyte's physiology, but the authoritative source is behind a paywall and unavailable. An AI assistant may supply a plausible equation — close to the real one, perhaps — without declaring that it is a reconstruction. This is hallucination in the precise sense: an approximation presented as if established.

The framework does not pretend to eliminate this. It reframes it. An undeclared approximation is an untracked reification: a pattern (the guessed equation) promoted into an object (a node of the reticulum) without provenance, on an unproven inferential step, with no validation. Controlled reification (§8) converts it into a typed artefact with open epistemic debt — missing provenance, unproven inference, empty validation state — and gates it: the object may be used provisionally but is blocked from authoritative downstream use until the debt is discharged against the real source. The question is thereby shifted from true versus false to declared versus undeclared epistemic state. An approximation with declared debt is honest and usable; the same approximation, silent, is the hazard. This is exactly the discipline that lets a human safely mine science with an AI that will sometimes, inevitably, fill a gap.

Mining science with AI: the undeclared approximation becomes declared debt



The framework does not eliminate hallucination. It converts an undeclared approximation into a typed object with declared epistemic debt — close to reality, perhaps, but never silently authoritative.

Figure 4. The undeclared approximation (hallucination) becomes a typed reification with declared, blocking debt — close to reality, perhaps, but never silently authoritative.

6. The Oracle: Navigating Governed Knowledge

The construction has a destination, and naming it clarifies the value of the whole programme. If knowledge is accumulated as a typed semantic reticulum with provenance, validation state and debt attached to every object, then a system can answer a query by navigating that reticulum rather than by regenerating an answer. This is a governed oracle: a central responder that traverses a typed semantic lattice and returns answers that are persistent and end-to-end traceable to their sources, that carry their declared debt and validation state, and that abstain or flag contradiction when the lattice cannot support a definite answer.

A concrete query makes the difference tangible. Asked “does this supplier-payment control cover every exception path?”, a fluent generator returns a confident yes or no. The oracle instead navigates the reticulum: it reaches the reified control object, finds the reification debt opened in §5 still undischarged on one exception branch, and — because that branch reduces to an undecidable coverage property (§9) — returns review-required with a trace to the unproven branch and the open debt, not a verdict. The answer is the path, and the not-knowing is explicit. That is the value proposition. A fluent generator produces an answer and discards the path; provenance and proof must be reconstructed afterwards, if at all. An oracle over a governed reticulum makes the path the product: every answer is reconstructable to sources; hallucinations surface as declared debt rather than silent error; knowledge composes in a governed substrate instead of being re-imagined per call; and not-knowing is a first-class, computable outcome. For high-stakes domains — science, compliance, legal, audit — this is a different paradigm: governed navigation of knowledge in place of ungoverned generation. The model ceases to be the knowledge and becomes its navigator. A-OSP (§11) is an early, honest attempt to build the substrate such an oracle requires.




Figure 5. From fluent generator to governed oracle: answers become persistent, traceable, debt-aware, and able to abstain.

BRIDGE — THE FRAMEWORK APPLIED TO ITSELF

6b. Self-Application: the Pipeline that Built the Witness

There is a piece of evidence that is neither pure theory nor an external study: the author applied the framework to itself. To build A-OSP, he instantiated the observer-compiler pattern as a pipeline of specialised AI agents — epistemic computational units in the sense of §7 — each producing a typed, validated intermediate artefact for the next. The pipeline has six stages: (1) capture the user intent, what is wanted inside A-OSP; (2) decompose that intent into engineerable blueprints; (3) check each blueprint against A-OSP's principles and prior decisions — the forbidden equivalences, the single source of truth, the proof chain; (4) decompose into homogeneous, mutually coherent sub-blueprints; (5) generate the roadmap as an agent-guided issue chain with explicit dependencies; (6) write complete issue bodies, cross-referenced, each carrying its definition of done, metrics and checklists, with coherence enforced both across issues and against the roadmap as a whole. Each stage is the corpus pattern in miniature — a prompt and prior artefacts produce a candidate, which review and audit close into a validated artefact — and the reification and coherence discipline (semantic CI/CD) is applied to the development artefacts themselves.

Self-application: the framework's agent pipeline built its own witness (A-OSP)




Figure 6. The six-stage agent pipeline by which the framework was applied to itself to construct A-OSP.

What this does and does not prove (efficiency caveat)

It does not claim A-OSP was built in the most efficient way, nor that A-OSP will be effective in itself. The hypothesis is asymmetric and stated openly: robustness and enterprise-MVP effectiveness are candidate — A-OSP may attain them; efficiency is, by design, not the goal.

Corollary, stated plainly: more efficient software suites may exist that do the same thing as A-OSP. The value claimed is governability, reconstructability and proof — not speed or economy. The self-application is therefore a witness of method feasibility — the framework is operable enough to have driven the construction of a non-trivial system — not a benchmark of effectiveness or efficiency.

PART II — THE GUARANTEES: WHY IT CAN BE TRUSTED

7. The Backbone, the Reuse It Declares, and the Positioning

Four layers carry the construction, read under the single thesis of computable semantics: RLA, the structural grammar of levels, transmissions and horizons (encoding included, since computability is encoding-sensitive); CRC, the conditions under which a reticulum is a computably operable unit relative to a declared horizon; ECNN, the operational mechanism of local-operator-plus-collapse over semantic fields, held as analogy, not a claim about neural substrates; and ROA, the governance compression of these into a pipeline. The constructs table below gives the per-layer detail; the figure shows the stack.

The reuse is declared, because declaring it is hygiene. The reticular grammar reorganises multi-level modelling; abstention and contradiction generalise selective prediction and the reject option; provenance discipline overlaps with established practice; the propagation result rests on Rice's theorem. None of these alone is the contribution. What is contributed is the layer that composes them into one governable observer — a contribution of integration, which described honestly is more defensible than a claim of novel invention, because it is precisely what the parts cannot do separately.

The four-layer backbone, read upward under one organizing thesis.




Figure 2. The four-layer backbone, read upward under one organizing thesis: RLA, CRC, ECNN, ROA.

Computable ≠ validated ≠ warranted

Four notions are kept distinct: computable (an effective procedure exists), epistemically computable (the system terminates with a structured epistemic state), validated (a domain procedure has accepted, rejected or revised the artefact), and warranted (it meets the domain standard of proof or evidence). This separation is what prevents the computable-semantics thesis from collapsing into a claim about truth. (Foundations: Annex A; ECNN/ECU: Annexes C, E.)

The positioning is best stated as differentiation, not criticism: each adjacent line of work is strong within its focus, and ROA adds a governance layer that none of them makes explicit. The table records, for each, its primary focus and what ROA adds.

Adjacent work	Primary focus	What ROA adds
Neural-symbolic AI	joint learning of representations and symbolic inference	pipeline-level governance, reification control and debt tracking around any such component
Selective prediction	when to abstain at the output	abstention/contradiction as typed objects that propagate across levels into review
Provenance / PROV	recording data and process lineage	proof vs projection vs witness vs approval as enforced, non-interchangeable transitions
Model/data documentation	describing models and datasets	operational enforcement at runtime rather than descriptive prose
KR / reification theory	warning against overcommitment	controlled reification as an auditable transition with quantified epistemic debt
RAG / agentic workflows	retrieval-conditioned generation; multi-step model + tool composition	retrieval and tool use as typed transmissions inside a reticulum; abstention and blocking debt as first-class terminal states

One linkage is what makes the contribution scientifically load-bearing rather than terminological: controlled reification and epistemic debt only mean something relative to the grammar of levels and transmissions. A reification is the promotion of a stabilised pattern to an object at the level above; debt is the residue of that move. Without levels and transmissions there is nowhere for either to live, and “reification” collapses into a generic warning about hypostatisation; with them, both become typed operations that can be audited, blocked or rolled back. The backbone is the precondition of the governance claim, not decoration.

For navigation, the principal constructs can be read at a glance, each living at exactly one backbone layer:

Layer	Construct	One-line meaning
RLA	level / transmission / horizon	states+rules+language+encoding; injective/quasi-/non-injective map; admissible questions and answer types
RLA	collapse / emergence	declared, scoped information loss; macro-property earns its keep only if it supports cheaper macro-regularities (Guard 3)
CRC	IOA / EC / TC; basic vs strong	compact, computably operable unit under a horizon; strong tier carries Turing-like structure and formal limits

ECNN	epistemic convolution / ECU	local-operator-plus-collapse over a semantic field; bounded transducer emitting typed artefacts (answer/unknown/contradiction)
ROA	reification chain / epistemic debt	recursive promotion of patterns into objects; residual obligation propagated until discharged or blocking
ROA	terminal states	answer / unknown / contradiction / out-of-horizon / review-required, all first-class

8. Controlled Reification as a Chain, and Epistemic Debt

Reification — the conversion of a relation, pattern or label into a manipulable object — is unavoidable and rarely happens once. It happens in stages: an interview answer is reified into atoms, atoms into links and evidence gaps, links into a synthesis, the synthesis into a reviewable artefact, the artefact into a governance record. Each stage produces the field of the next, and an unvalidated promotion at one stage becomes the raw material of the next — so any object built on it inherits the unresolved obligation. Controlled reification makes each move auditable, requiring source references, scope, validation state, debt metadata, allowed uses and a rollback path.

Epistemic debt is that residual obligation — to be made visible, assignable, reviewable, and where appropriate blocking on downstream use, not necessarily eliminated. The reification component is operationalised: the reification-debt rate over a session is the count of reified objects consumed downstream with empty or non-authoritative validation state, over total reified objects consumed downstream. Controlled reification drives this rate towards zero; any non-zero value localises exactly where an unvalidated construct leaked into a decision. This makes the central contribution falsifiable rather than descriptive — the same mechanism that, in §5, turned the hallucinated equation into declared, blocking debt.

Running example. In the supplier-payment review, interview transcripts (bounded domain material) collapse non-injectively into atoms — many phrasings map onto the single atom “a CFO-approval threshold exists” — while the meaning-critical distinction between a control that exists and one whose evidence is missing must survive the collapse. Promoting “control_gap” into a reusable object opens reification debt; the object is usable for drafting but blocked from the governance record until a reviewer validates it. A guarded note on collapse: the framework does not assert as an axiom that every non-injective collapse yields a genuine macro-property — that reading is near-analytic. A macro-property counts as reticularly emergent only if it is predictively autonomous: there are macro-level regularities in which it figures and which are cheaper to compute at the upper level than by reconstructing the lower state. Collapse alone is bookkeeping; an emergent property must earn its keep.

9. Propagation of Limits: a Reductio, and Its Governance Twin

Here the heavy theory does real work, by defeating a tempting intuition: compositional decidability, the belief that a system assembled from locally computable levels must itself be tractable. It is false. A composite of locally computable rules — a layered pipeline, or in the simplest illustration a videogame, which is nothing but locally computable rules stacked into a composite — can have extensional properties that are undecidable, though no single rule is. Undecidability is a property of the composition, not of the pieces.

The propagation argument is a reductio, not an assumption

Let level L_i be Turing-like on a critical set C_i , and let the transmission to L_{i+1} be effectively computable and injective on C_i . Suppose a non-trivial extensional property, undecidable at L_i by Rice (1953), induced a decidable property at L_{i+1} . Then composing the transmission with the hypothetical decider would decide the property at L_i — contradicting Rice. Hence undecidability propagates forward wherever critical distinctions are preserved; where a transmission feeds a higher level back into a critical level, the limit reverberates backward.


Status: the reductio is sound in outline. What remains owed is the complete statement under minimal hypotheses — the quasi-injective case and the feedback dynamics — listed as a research obligation (§13). The gap is between a sound argument and a fully pinned theorem, not between an assumption and an argument. (Foundations and CA bridge: Annexes A, D.)

From the propagation statement follows mandatory abstention: if a workflow question reduces to a non-trivial extensional property of the behaviour computed at a Turing-like level, no downstream tooling can convert it into a decidable query, and the only sound terminal state is structured abstention. The strongest theoretical claim is the duality. Undecidability propagates through distinction-preserving transmissions; in governance settings, unresolved obligations propagate through any downstream transition that reuses an unsupported distinction, forcing blocking debt. Non-injective collapse may block the first kind of propagation but creates the second whenever the collapsed distinctions remain relevant downstream. The two regimes share a structural shape and differ in remedy: mandatory abstention is structural and irreducible; blocking debt is procedural and dischargeable through review, evidence or escalation.


Proposition 1 (mandatory abstention) and the running example

If a workflow question q at a Turing-like level reduces to a non-trivial extensional property U of the behaviour computed there, then U is undecidable (Rice, 1953), and no retrieval, model call, governance workflow or interface state can make q decidable. The only sound terminal state is structured abstention; a system that returns a definite answer to q is not miscalibrated but unsound.

In the supplier-payment review, “do the documented controls cover every possible exception path?”, posed over the control procedures treated as a program, is such a property: it cannot be settled by more interviews or more model calls, and its correct terminal state is horizon-exceeded. By contrast “is there documented evidence of a CFO-approval threshold?” is bounded retrieval and is answerable. CRC-strong is what tells the architecture, in advance, which question is which.

Formal track — undecidability

with feedback: reverberates backward

Governance track — epistemic debt

Undecidability propagates when critical distinctions are preserved; debt propagates when unsupported distinctions are reused. Abstention is structural and irreducible; blocking debt is procedural and dischargeable.

Figure 3. The propagation duality. Undecidability propagates when critical distinctions are preserved (and reverberates backward under feedback); debt propagates when unsupported distinctions are reused.

10. Normative Meta-Control

In normative domains the architecture adds meta-control over normative boundaries, modelled over a world model (facts, evidence, obligations), a self model (limits, calibration, drift, known gaps) and a normative kernel (rules, policies, constraints), with an epistemic history that makes the process path-dependent. The meta-controller consumes the governance signals the reification chain and debt propagation produce: it is where unsupported claims, missing evidence, modal drift and unresolved blocking debt are escalated to human authority rather than silently accepted. It does not decide moral truth and does not authorise itself — its function is to make the boundary between the system's reach and human authority visible, not to move it. Read operationally, this is the sober core of what the broader work calls proto-consciousness: epistemic self-boundedness, the capacity to compute when a query crosses what the observer's ontology and horizon can support, and to emit structured non-knowledge. (Meta-controller: Annex G; the speculative topological reading is developed only in the main paper and Annex D and is not relied upon here.)

11. A-OSP: Reference Architecture and the Forbidden Equivalences

A-OSP is a reference architecture: an implementation witness for how the framework can be expressed as a runtime substrate, and the early scaffolding of the oracle of §6. It is not empirical validation — framework and architecture co-evolved, so it is an existence proof that the abstraction is implementable, not evidence that it is beneficial. It is browser-native and text-first, with a durable source of truth in typed, append-only records, and keeps six object classes non-interchangeable — source, cache, view, adapter, export and proof. A semantic CI/CD layer tests meaning and coherence (blueprint alignment, issue consistency, naming drift, unresolved assumptions, reification leaks): the same discipline the self-application pipeline (§6b) applied to its own artefacts.

The architectural thesis is that proof is a chain, not a badge: a proof-sensitive transition becomes proof-grade only when the object is persisted, read back, compared, represented by a receipt and exposed as proof state — distinguishing a genuine proof from a false green. The most transferable element is a set of enforceable non-equivalences that translate directly into negative tests; they are the engineering form of the contribution.

Forbidden equivalences (enforceable, testable)

output $\neq$ proof · confidence $\neq$ evidence · log $\neq$ receipt · export $\neq$ witness · interface state $\neq$ proof · model memory $\neq$ source of truth · review $\neq$ approval · generated document $\neq$ authority source

A-OSP is built to learn, not to certify, and its progress is gated by one deliberately narrow proxy: proof-grade D1 (intake) and EQL (scoped retrieval) — receipt-backed, read-back-confirmed, resistant to false-green interface states. Its roadmap is credible in a limited sense: all fourteen identified proof-grade gaps are owned by open issues, which makes the plan honest without converting planned work into completed proof. Closing the proxy yields the first datum on whether the typed-object discipline reduces downstream error, and is the smallest honest step from oracle-as-vision to oracle-as-evidence. Per-surface maturity is declared, not assumed (AS-IS / HARDENING / TARGET / HYPOTHESIS / NON-GOAL).

12. Claim Discipline and Scope

Constructs are classified by burden of proof, so that theses, contributions, conjectures and research obligations are not judged by one evidential standard.

Construct	Status	Safe formulation
Computable semantics	Organizing thesis	observation and stabilization of meaning is an effective, bounded, horizon-relative process emitting typed artefacts; not a claim about truth
Science as node / epistemic twin	Method	a scientific domain is mined into a compact computable reticulum that is the observer's horizon and source of truth; demonstrated on the bryophyte
Epistemic indistinguishability	Falsifiable prediction	blind experts cannot separate synthetic from empirical trajectories above a threshold; design goal, not validated
Hallucination as declared debt	Central practical claim	an undeclared approximation is an untracked reification; typed, given open debt, blocked from authoritative use until discharged
Governed oracle	Vision + early witness	answers by navigating a typed semantic reticulum; persistent, traceable, debt-aware; substrate prototyped in A-OSP, proxy not yet closed
Self-application pipeline	Method-feasibility witness	framework used to build its own witness via a 6-stage agent pipeline; not a claim of efficiency or effectiveness; more efficient suites may exist
Reification chain + debt propagation	Central contribution	recursive auditable promotion, debt propagated until discharged or blocked; D_reif(rate) operationalised
Propagation of undecidability	Reductio; hypotheses owed	sound in outline by reduction to Rice; complete statement under minimal hypotheses on quasi-injective restriction and feedback is owed
Mandatory abstention vs blocking debt	Terminal regimes	abstention structural and irreducible; blocking debt procedural and dischargeable

Construct	Status	Safe formulation
Epistemic convolution / ECU	Design abstraction	CNN-inspired local-operator pattern over semantic fields, held as analogy; ECU a bounded transducer under matrix and envelope
Proto-consciousness (operational)	Sober reading kept	epistemic self-boundedness: computing when a query exceeds the supportable horizon; the speculative topological version is excluded here
A-OSP	Reference architecture	implementation witness, co-evolved; success proxy = proof-grade D1/EQL; not independent validation

13. Open Problems and Falsification

The obligations are stated, because stating them is part of the discipline. The propagation reductio (§9) must be completed into a theorem under explicit minimal hypotheses, including the quasi-injective case and feedback dynamics. Epistemic debt is operationalised here only for the reification component and only as a rate; the other components need their own protocols with baselines. The bryophyte’s epistemic indistinguishability (§4) is a design goal awaiting a blind expert study. The complementarity with neural-symbolic systems is argued, not demonstrated: a reproducible comparison of a reticular observer instantiation, an unstructured baseline and a neural-symbolic system on a bounded compliance task — reporting accuracy, abstention, contradiction detection, indecidability recall, provenance preservation, reification-debt rate and review effort, under a Popper-style challenge harness — is owed. A-OSP is a witness, not validated evidence; closing the D1/EQL proxy is the first empirical test and the smallest step toward the oracle.

Falsification of the framework

The framework would be shown unhelpful — as a framework, not merely in one system — if systems built under it showed no reduction in downstream decision error relative to unstructured baselines on matched tasks; if the overhead of typed artefacts, receipts and reification control consistently exceeded the governance benefit; if the reification-debt rate failed to correlate with any independent measure of downstream error; or if, on the bryophyte, experts reliably separated synthetic from empirical trajectories at the chosen observables. These are measurable, and they are the conditions under which the framework should be abandoned or substantially revised.

14. Conclusion

AI systems should not be evaluated only as answer generators. The constructive claim is that, once semantics is treated as computable in an observer-relative sense, one can mine a scientific domain into a typed computable reticulum — an epistemic twin governed by its own horizon — and answer from it as a traceable, debt-aware oracle; the author has gone as far as applying the framework to itself to build the substrate, with the explicit caveat that this witnesses feasibility, not efficiency. The disciplinary claim is that the same construction makes the characteristic failure of AI visible and governable: hallucination becomes typed reification with blocking debt, and not-knowing becomes first-class. The heavy theory earns its place by defeating compositional-decidability intuition and revealing the governance twin of undecidability in blocking debt. The remaining obligations — a pinned theorem, a blind bryophyte study, a comparative evaluation, a closed proof proxy — are stated, local and measurable. The future of trustworthy AI-assisted work will depend less on more fluent generation than on making the observers behind generation explicit, computable and governable.

Methodological Note and Relation to the Underlying Work

The conceptual development, formal definitions and claim classification were directed by the human author, who retains full scientific responsibility. AI assistants supported literature triage, structural revision, consistency checks and — as described in §6b — the staged construction of the reference architecture’s blueprints and issues; all references were manually checked. This paper is the upper, self-contained layer of a longer development by the same author: beneath it sit an integrative methodological paper (“AI as Observer Compiler”), an engineering whitepaper for the reference architecture (A-OSP), a chapter-aligned reply to Wolfram’s metaphysics — in which laws and objectivity are read as stability classes induced by admissible compressions, rulial transduction is domesticated into governed practice, and global necessity claims are admitted only when an identification procedure remains available — and a technical corpus: foundations of RLA and CRC (Annex A), the ECNN formalisation (Annex C), the PCE/cellular-automata bridge (Annex D), the ECU/UCE specification (Annex E), the bryophyte CRC case (Annex B), prototype implementations including a textual ECU oracle (Annex F), and the normative meta-controller (Annex G). These are noted for transparency; no claim in this paper depends on access to them.

Corpus Navigation Addendum

This addendum makes the whole corpus navigable from ROA v6. It lists every document of the public repository, explains each in one paragraph, gives reading paths by need, and provides two cross-reference tables: a formal one (where each construct is canonically defined and where ROA v6 uses it) and a notational one (each symbol → source → ROA v6 section). ROA v6 remains self-contained; this addendum is for readers who want to descend from it into the technical layer. Note on completeness and honesty: all documents below were read directly except the PCE bridge slide deck, which is described from the repository listing as the presentation companion to the main paper and Annex D, not from independent reading.

1. Corpus hierarchy

```

Reticular Observer Architectures for Governable AI-Assisted Work [ENTRYPOINT]
└ (i) AI as Observer Compiler — from Wolfram's Ruliad to RLA-ECNN
└ (ii) Augmented Ontological Semantic Platform (A-OSP) Whitepaper
|
└ RLA-CRC-ECNN/
|   └ Main Paper RLA-ECNN-CRC-PCE      [CORE THEORY]
|   └ Slide deck RLA-ECNN bridge PCE    [presentation companion]
|   └ annex A RLA-CRC foundations
|   └ annex B RLA biological case (bryophyte)
|   └ annex C ECNN formalisation
|   └ annex D RLA-ECNN / PCE bridge
|   └ annex E ECU-UCE specification
|   └ annex F proto-epistemic architectures (prototypes)
|   └ annex G methodology / experiments (iKant meta-controller)
|
└ Reply to Wolfram/
  └ AI as Observer Compiler — ANNEX — reply to Wolfram's metaphysics via RLA-ECNN

```


2. Navigation table

For each document: its role, what it canonically contributes, and the ROA v6 sections it grounds or is grounded by.

Document	Role	Canonical contribution	ROA v6 link
ROA (entrypoint)	Governance compression; standalone thesis	controlled reification + epistemic debt propagation as first-class transitions; the two-lung argument	is ROA v6
(i) AI as Observer Compiler	Integrative methodology	observer-compiler view; intermediate artefact chains; prompt-as-operator; knowledge mining; rulial transduction→domestication	§1, §3, §6b, Note
(ii) A-OSP Whitepaper	Engineering / implementation witness	filesystem-first SSOT; six object classes; forbidden equivalences; proof chain; D1/EQL proxy; semantic CI/CD	§6, §6b, §11
Main Paper RLA-ECNN-CRC-PCE	Core theory position paper	unifies RLA/CRC/ECNN; propagation + emergence; bryophyte summary; proto-consciousness (speculative)	§2, §7, §9, §11
Slide deck (PCE bridge)	Presentation companion	visual summary of the PCE bridge (see Annex D); not independently load-bearing	context for Note
Annex A — foundations	Formal foundations of RLA/CRC	definitions/axioms: $R=(L,T,H)$, τ classes, IOA/EC/TC, $CR(R,H)$, collapse/emergence, propagation	§7, §8, §9
Annex B — bryophyte	Worked scientific twin	compact RLA topology (L01–L20, P01–P42); epistemic indistinguishability (blind Turing test)	§3, §4
Annex C — ECNN	ECNN formalisation	convolution/pooling as RLA transmissions; epistemic head emitting unknown/contradiction	§7
Annex D — PCE bridge	Bridge to Wolfram's PCE / CA	RLA-CA mapping; propagation alignment with Annex A; falsifiable conjectures	§9, Note
Annex E — ECU/UCE	Unit specification + design patterns	ECU ϕ : $Z \times M \rightarrow A$; matrix $M=(E,R,C,P)$; functional rings; LLM-as-constrained-transducer	§7

Annex F — prototypes	Prototype implementations	bryophyte simulator; ECNN benchmark; textual ECU oracle; UCE; RLA-CA; Popper- χ harness	§6, §13
Annex G — methodology (iKant)	Normative meta-controller	world/self/norm models; epistemic history; case- and system-level outputs; escalation	§10
Reply to Wolfram	Chapter-aligned metaphysics reply	laws/objectivity as stability classes under admissible compressions; demarcation rule for global necessity; proto-consciousness as self-boundedness	§10, Note

3. Per-document local explanations

ROA — entrypoint (this corpus)

The governance compression and standalone thesis: once semantics is treated as computable in an observer-relative sense, a scientific domain can be mined into a typed reticulum and answered as a governed oracle, while the characteristic failure of AI — the undeclared approximation — is converted into typed reification with blocking debt. Defends one contribution: controlled reification and epistemic debt propagation as first-class, typed, auditable transitions. Refers down to every document below.

(i) AI as Observer Compiler

The integrative methodology beneath ROA. It reframes prompt→answer as the construction of intermediate, typed artefacts (definitions, blueprints, issue chains, checklists) curated by a human across calls — the basis of the self-application pipeline in ROA v6 §6b. Introduces knowledge mining, original epistemic artefact synthesis, prompt-as-operator, semantic debt, and the rulial transduction→domestication framing connecting Wolfram's Ruliad to bounded reticular computation.

(ii) A-OSP Whitepaper

The engineering whitepaper for the reference architecture: a browser-native, text-first epistemic operating environment whose durable source of truth is typed, append-only records. It specifies the six non-interchangeable object classes, the forbidden equivalences, the proof chain (persist→read-back→compare→receipt→proof state), semantic CI/CD, and the narrow D1/EQL proof-grade success proxy. It is an implementation witness, not validation, and is the substrate of the oracle of §6.

Main Paper RLA-ECNN-CRC-PCE

The core theory position paper that unifies the layers and states the propagation of undecidability and the guarded emergence relation, summarises the bryophyte case, and outlines — explicitly as speculation — the topological proto-consciousness conjecture and the PCE bridge. ROA v6 draws its backbone (§7), propagation (§9) and meta-control (§11) from here, while excluding the speculative extensions from its claims.

Slide deck (RLA-ECNN bridge PCE)

The presentation companion to the PCE bridge material (Annex D and the main paper). Described here from the repository listing rather than independent reading; it carries no claim not already in the main paper or Annex D, and is included for completeness of the corpus map.

Annex A — RLA/CRC foundations

The formal foundations: reticulum $R=(L,T,H)$; level $Li=(Di,\Sigma,Langi,Enci)$; transmission τ with injective / quasi-injective / non-injective classes; collapse coefficient CC and indices IE, RIR ; the triple-guarded reticular emergence definition; $IOA/EC/TC$ and $CR(R,H)$; CRC-basic and CRC-strong; and the Rice-style propagation statement. This is where the symbols ROA v6 uses are canonically defined.

Annex B — bryophyte case

The worked epistemic twin: a fully computable RLA topology for a generalist bryophyte (twenty levels $L01-L20$, forty-two parameters $P01-P42$ with literature ranges, update equations, collapse/emergence metrics, climate scenarios). It defines epistemic indistinguishability and the blind expert test (a Turing test for computed versus measured science) that ROA v6 §4 presents as a falsifiable prediction.

Annex C — ECNN formalisation

The formalisation of Epistemic Convolutional Neural Networks: convolutional and pooling layers reread as RLA levels with non-injective coarse-grainings, plus an epistemic head and ECU that can accept, reject, or emit unknown/contradiction while remaining Turing-computable. Grounds ROA v6 §7's treatment of epistemic convolution as a design analogy.

Annex D — PCE bridge

The bridge layer aligning RLA/CRC and ECNN with Wolfram's Principle of Computational Equivalence via one-dimensional cellular automata, with the propagation reading consistent with Annex A and falsifiable conjectures rather than ontological claims. Supports ROA v6 §9 and the methodological note; the bridge itself is not load-bearing for any ROA v6 claim.

Annex E — ECU/UCE specification

The unit specification: an ECU ϕ mapping a representation Z and an epistemic matrix $M=(E,R,C,P)$ to a structured artefact carrying status, rationale, debt and validation state; higher-order reticula of ECUs as functional rings (executive, valuation, memory, metacognition); and the regime in which an LLM is a constrained transducer under an operational envelope, not an oracle. Grounds ROA v6 §7.

Annex F — proto-epistemic architectures

Prototype implementations and experimental setups: the bryophyte CRC simulator, an ECNN image-classification benchmark with epistemic decisions, textual ECUs as epistemic oracles (the seed of ROA v6 §6's oracle), a UCE environment, an RLA-CA explorer, and a generic Popper- χ challenge harness. Supports ROA v6 §6 and the comparative study owed in §13.

Annex G — methodology / iKant

The normative meta-controller formalised as the iKant pattern: a triple decomposition into world model, self model and normative kernel, with an epistemic history giving path-dependence, producing case-level and system-level outputs and escalation. Grounds ROA v6 §10; iKant is a governance pattern, not moral agency.

Reply to Wolfram

A chapter-aligned reply to Wolfram's metaphysics of the Ruliad, using epistemic horizons, multi-level transmissions and compact reticular computability to make agreement and disagreement checkable. It reads laws and objectivity as stability classes induced by admissible compressions, admits global “must exist” claims only when an identification procedure remains available, and gives the operational reading of proto-consciousness as epistemic self-boundedness that ROA v6 §10 adopts.

4. Formal cross-reference (where each construct is defined and used)

Read: construct → the document where it is canonically defined → the ROA v6 section that uses it.

This lets a reader descend from any ROA v6 claim to its formal home.

Construct	Canonical definition	Use in ROA v6
Reticulum $R=(L,T,H)$; level $L_i=(D_i,\Sigma_i,Lang_i,Enc_i)$	Annex A	§7 (backbone), §8
Transmission τ (injective / quasi-injective / non- injective)	Annex A	§7, §9
Collapse coefficient CC; indices IE, RIR	Annex A	§8 (guarded note)
Reticular emergence (Guards 1–3; computational mechanics)	Annex A; Main	§8
IOA, EC, TC; $CR(R,H)$; CRC-basic / CRC-strong	Annex A; Main	§7, §9
computable $\neq$ validated $\neq$ warranted	Main; ROA	§7 (callout)
Rice-style propagation of undecidability	Annex A; Annex D	§9 (reductio)
Proposition 1 (mandatory abstention)	Main; ROA	§9 (callout)
Propagation duality (undecidability $\Leftrightarrow$ debt)	ROA	§9
Epistemic convolution layer; pooling as collapse	Annex C	§7
ECU $\phi: Z \times M \rightarrow A$; matrix $M=(E,R,C,P)$; functional rings	Annex E; Annex C	§7
Epistemic artefact fields (status / rationale / debt / validation_state)	Annex E	§6, §8
Controlled reification; reification chain	ROA; Observer Compiler	§5, §8
Epistemic debt; $D_reif(rate)$; 9 components	ROA; Observer Compiler	§8

Forbidden equivalences; six object classes	A-OSP	§11, §12
Proof chain; proof-grade; D1 / EQL; false-green	A-OSP	§11
Semantic CI/CD	A-OSP; Observer Compiler	§6b, §11
Intermediate artefact chains; prompt-as- operator	Observer Compiler	§1, §6b
iKant meta-controller (world / self / norm; epistemic history)	Annex G	§10
Proto-consciousness: operational (self- boundedness)	Reply to Wolfram; Main	§10
Proto-consciousness: topological (speculative)	Main; Annex D	excluded (cross-ref only)
Rulial transduction → domestication	Observer Compiler; Reply to Wolfram	Note
Epistemic indistinguishability (blind test, threshold η)	Annex B	§4
Bryophyte topology (L01– L20, P01–P42)	Annex B	§4
Popper-χ harness; metrics (CR, IR, ...)	Annex F; Main	§13

5. Notation cross-reference (symbol map)

A compact symbol table, so any notation met in ROA v6 or the annexes resolves to one meaning and one canonical source.

Symbol	Meaning	Defined in
R = (L, T, H)	reticulum: levels, transmissions, horizon	Annex A
Li = (Di, Si, Langi, Enci)	level: states, local rules, language, encoding	Annex A
$\tau \rightarrow j$	transmission (function or relation) between levels	Annex A
Ci	critical subset on which τ is (quasi-)injective	Annex A; Annex D
$CC(\tau, S) = 1 - \tau(S) / S$	collapse coefficient over a finite set S	Annex A

IE, RIR	emergence index; reticular irreducibility index	Annex A
IOA, EC, TC	internal ontological autonomy; epistemic closure; Turing-computability	Annex A
CR(R,H)	compactness predicate = $IOA \wedge EC \wedge TC$ under horizon H	Annex A
H	epistemic horizon (admissible questions, sources, ops, answer types)	Annex A; ROA
U	non-trivial extensional property (undecidable at a Turing-like level)	Annex A; Main
$\phi: Z \times M \rightarrow A$	ECU map: representation $\times$ epistemic matrix $\rightarrow$ artefact space	Annex E
$M = (E, R, C, P)$	epistemic matrix: entities, rules, coherence, policies	Annex E; Annex C
ρ (envelope)	operational envelope of an LLM-as-ECU (snapshot, decoding, schema, parser)	Annex E; Observer Compiler
$hk+1(v) = \text{Aggk}(\{ECU_k, m(\cdot)\})$	epistemic convolution layer aggregating ECU outputs over a neighbourhood	Annex C
D_reif (rate)	reified objects used downstream with empty validation / total	ROA
A_cand, A_inter, A_val	candidate / intermediate / validated artefacts in the chain	Observer Compiler
W, S, N	world model, self model, normative kernel (iKant)	Annex G
η	discrimination threshold in the blind indistinguishability test	Annex B
$\Omega_{\text{Ruliad}} \rightarrow R_{\text{local}}$	rulial transduction to a local observer reticulum	Observer Compiler; Reply

6. Reading paths by need

If you want...	Start here	Then read
Fast orientation	ROA v6 (entrypoint)	this addendum, then the README
Core theory	Annex A	Main paper, Annex C, Annex E
AI / ML architecture	Annex C	Annex E, Annex F, ROA v6 §7
Scientific-modelling case	Annex B	Annex A, Main paper, ROA v6 §4

Experiments / prototypes	Annex F	Annex C, Annex E, ROA v6 §6,§13
Governance / compliance	ROA v6	Annex G, A-OSP whitepaper
Wolfram / Ruliad / PCE	AI as Observer Compiler	Reply to Wolfram, Annex D, slide deck
Implementation architecture	A-OSP whitepaper	ROA v6 §6b,§11, Annex F, Annex G

The corpus does not claim completed mathematical proof, empirical validation, production readiness, legal certification, or artificial consciousness — see ROA v6 §12–§13 for the claim discipline and the open obligations.

References

- Anderson, P. W. (1972). More is different. *Science*, 177(4047), 393–396.
- Besold, T. R., d’Avila Garcez, A., Bader, S., et al. (2017). Neural-symbolic learning and reasoning: A survey and interpretation. *arXiv:1711.03902*.
- Chow, C. K. (1970). On optimum recognition error and reject tradeoff. *IEEE Trans. Information Theory*, 16(1), 41–46.
- Church, A. (1936). An unsolvable problem of elementary number theory. *American Journal of Mathematics*, 58(2), 345–363.
- Geifman, Y., & El-Yaniv, R. (2017). Selective classification for deep neural networks. In *NeurIPS*.
- Glime, J. M. (2007). *Bryophyte Ecology, Vol. 1: Physiological Ecology*. Michigan Technological University.
- Godfrey-Smith, P. (2007). Conditions for evolution by natural selection. *Journal of Philosophy*, 104(10), 489–516.
- Humphreys, P. (2016). *Emergence: A Philosophical Account*. Oxford University Press.
- Kitano, H. (2002). Systems biology: A brief overview. *Science*, 295(5560), 1662–1664.
- Manhaeve, R., Dumančić, S., Kimmig, A., Demeester, T., & De Raedt, L. (2018). DeepProbLog. In *NeurIPS*.
- Mitchell, M., Wu, S., Zaldivar, A., et al. (2019). Model cards for model reporting. In *Proceedings of FAT**.
- Noble, D. (2012). A theory of biological relativity: No privileged level of causation. *Interface Focus*, 2(1), 55–64.
- Popper, K. (1959). *The Logic of Scientific Discovery*. Hutchinson.
- Proctor, M. C. F. (2000). The bryophyte paradox: Tolerance of desiccation, evasion of drought. *Plant Ecology*, 151(1), 41–49.
- Rice, H. G. (1953). Classes of recursively enumerable sets and their decision problems. *Transactions of the AMS*, 74(2), 358–366.
- Serafini, L., & d’Avila Garcez, A. (2016). Logic tensor networks. *arXiv:1606.04422*.
- Shalizi, C. R., & Crutchfield, J. P. (2001). Computational mechanics. *Journal of Statistical Physics*, 104(3–4), 817–879.
- Turing, A. M. (1937). On computable numbers, with an application to the Entscheidungsproblem. *Proc. London Math. Soc.*, s2-42(1), 230–265.
- Vovk, V., Gammerman, A., & Shafer, G. (2005). *Algorithmic Learning in a Random World*. Springer.
- W3C. (2013). *PROV-O: The PROV Ontology*. World Wide Web Consortium Recommendation.
- Wolfram, S. (2021). *What Ultimately Is There? Metaphysics and the Ruliad*. Wolfram Media.